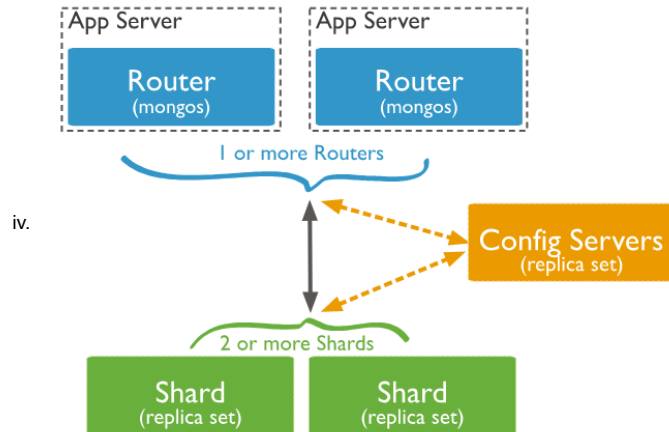


# Clustering

Tuesday, October 28, 2025 4:05 AM

## 1. Clustering

- a. Always a 3 shared cluster (3 minimum)
  - i. **Shard:** Each shard contains a subset of the sharded data. Each shard must be deployed as a replica set.
  - ii. **Routing with mongos:** The mongos acts as a query router, providing an interface between client applications and the sharded cluster.
  - iii. **Config servers:** Config servers store metadata and configuration settings for the cluster. Config servers must be deployed as a replica set (CSRS).
- b. Why clustering
  - i. To avoid single point of failure.
  - ii. To boost performance while querying.
- c. Sharding
  - i. Using shard keys: consist of a field or multiple fields in documents
  - ii. Missing shard key fields are treated as having null values
    1. Reshard a collection
  - iii. Shard key Index
    1. Choice of shard key affects performance and scalability of the cluster
    2. Shard key index plays a crucial role in querying
    3. Refer to this guide:
      - a. <https://www.mongodb.com/docs/manual/core/sharding-choose-a-shard-key/#std-label-sharding-shard-key-selection>



- v. Refer to the sharding guide:
  1. <https://www.mongodb.com/docs/manual/sharding/>
- d. Advantages of Sharding
  - i. MongoDB distributes the R/W ops across the shards in the cluster
    1. This allows them to process a subset of cluster operations
    2. Both R/W workloads can be scaled horizontally across the cluster.
      - a. Targeted ops more efficient than broadcasting to every shard
    3. Increase storage capacity
    4. High Availability

## 2. Streaming Data

- a. Integrations with Apache Kafka
- b. Respond to real time events with triggers