

Transactions & Atomicity

Tuesday, October 28, 2025 2:09 AM

1 Transactions in MongoDB

MongoDB supports **multi-document ACID transactions** since **v4.0** (for replica sets) and **v4.2** (for sharded clusters).

Key Points:

- **Single-document operations** are always atomic.
- **Multi-document transactions** ensure **ACID guarantees**: Atomicity, Consistency, Isolation, Durability.
- Transactions are created using **sessions** in MongoDB drivers:

Example (Java/MongoDB):

```
ClientSession session = mongoClient.startSession();
try {
    session.startTransaction();

    collection1.insertOne(session, doc1);
    collection2.updateOne(session, filter, update);

    session.commitTransaction();
} catch (Exception e) {
    session.abortTransaction();
} finally {
    session.close();
}
```

- **Isolation:** MongoDB transactions provide **snapshot isolation**, meaning each transaction sees a consistent snapshot of the data at its start.

2 Handling Concurrent Requests (Two Users Hitting Same Request)

MongoDB uses **document-level concurrency**:

- Only one operation can modify a single document at a time.
- If two users try to update the **same document**, the **first write wins**, the second write may fail if you rely on conditional updates (updateOne with a filter on a version or timestamp).

Patterns for handling concurrency:

a) Optimistic Concurrency Control (OCC)

- Add a version field to the document.
 - Update only if version matches:
- ```
collection.updateOne(
 Filters.and(Filters.eq("_id", docId), Filters.eq("version",
currentVersion)),
 Updates.combine(
 Updates.set("field", newValue),
 Updates.inc("version", 1)
)
);
```

- If update count = 0 → retry or abort.

#### b) Pessimistic Locking (Limited in MongoDB)

- MongoDB **does not support row-level locks like SQL**.
- You can emulate locks with a **"lock" document**:  

```
{ "_id": "resource_lock", "lockedBy": "instance1", "timestamp": 12345678 }
```
- Use findOneAndUpdate with a filter to acquire lock atomically.
- Release after operation.

## 3 Multiple Instances Writing to DB

When **multiple app instances** write:

- MongoDB **replica sets handle replication**; writes go to the primary.
- To prevent **race conditions**, use **atomic operations** (\$set, \$inc, \$push, findOneAndUpdate).
- For **distributed locks**, you can implement a **lock collection** (like Redis locks) or use MongoDB's findOneAndUpdate with a TTL index for auto-expiry.

## 4 Comparison with Spring Transactions + ShedLock

| Aspect                   | MongoDB                            | Spring + RDB + ShedLock     |
|--------------------------|------------------------------------|-----------------------------|
| <b>Transaction Scope</b> | Session-based; multi-document ACID | Method-level @Transactional |

|                                 |                                                    |                                                                                 |
|---------------------------------|----------------------------------------------------|---------------------------------------------------------------------------------|
| <b>Isolation</b>                | Snapshot isolation                                 | Read Committed / Repeatable Read / Serializable                                 |
| <b>Concurrent Writes</b>        | Use OCC / locks at document level                  | Database handles row-level locks; @Transactional ensures atomicity              |
| <b>Distributed Jobs / Locks</b> | Implement manual lock collection or TTL-based lock | ShedLock library ensures only <b>one node runs a scheduled job</b> in a cluster |
| <b>Locking Mechanism</b>        | Atomic updates, optionally explicit locks          | Database row-level locking or explicit ShedLock table row                       |

**Key takeaway:**

- **MongoDB relies more on atomic operations and optimistic concurrency**, while **Spring + RDB** relies on DB locks and transaction isolation.
- For distributed systems with multiple nodes, **ShedLock** is a neat abstraction; in MongoDB, you implement a **lock collection** or use Redis/Zookeeper to prevent double execution.

**5 Example: Distributed Lock in MongoDB**

```
Document lockDoc = new Document("_id", "job1")
 .append("lockedAt", Instant.now())
 .append("lockedBy", instanceId);

Document result = lockCollection.findOneAndUpdate(
 Filters.eq("_id", "job1"),
 new Document("$setOnInsert", lockDoc),
 new
FindOneAndUpdateOptions().upsert(true).returnDocument(ReturnDocument.AFTER)
);

if (result.getString("lockedBy").equals(instanceId)) {
 // acquired lock, safe to run job
}
```

- **\$setOnInsert + upsert** ensures only one instance acquires lock.
- Release lock after job completes by deleting the document.

**✅ Summary / Practical Advice**

1. **MongoDB transactions** are ACID but scoped by session; prefer **atomic single-document updates** where possible.
2. **Concurrent writes** → use **optimistic concurrency** or atomic operations.
3. **Multiple instances** → handle with **atomic updates + lock collection** for distributed locks.
4. **Spring analogy**: transactions are like Mongo sessions; ShedLock is like a distributed lock collection in MongoDB.

